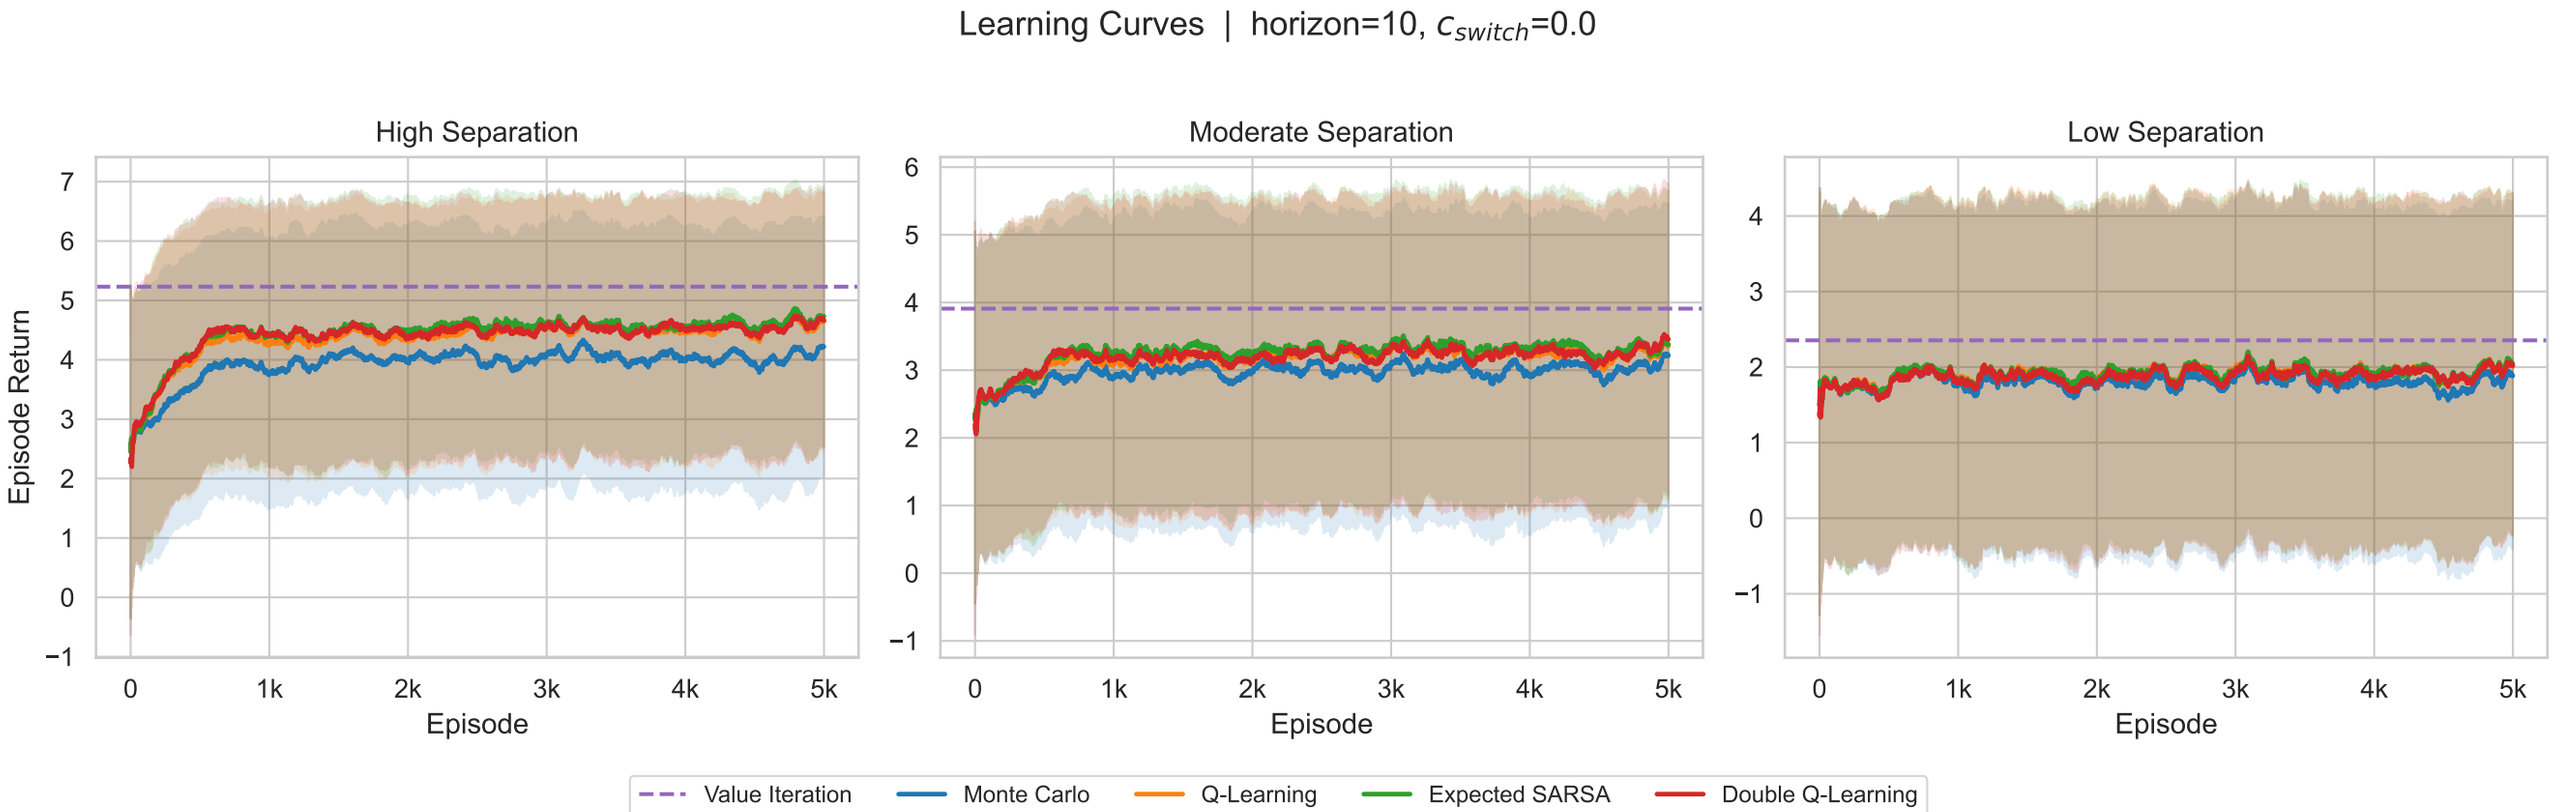
Learning Curves | horizon=10, c_{switch}=0.0
High Separation
Moderate Separation
Low Separation
Episode Return
Episode
0
1k
2k
3k
4k
5k
Value Iteration
Monte Carlo
Q-Learning
Expected SARSA
Double Q-Learning